

A dark blue vertical bar runs down the left side of the page. A blue arrow points to the right from this bar, containing the text 'Big Data Lab'. In the bottom left corner, there are several thin, curved, light blue lines that sweep upwards and to the right.

Big Data Lab

Hive Installation

Hadoop Setup For Hive

- Check your java version, if it is not Java8, then you need to install Java8 and set it as the default version.

```
java -version
```

- Steps to install java 8

```
sudo apt update  
sudo apt install openjdk-8-jdk -y
```

- Select java 8 version as the default version from the existing versions.

```
sudo update-alternatives --config java  
(Type the selection number corresponding to java 8)
```

- -Create a java.sh file and export the JAVA_HOME and PATH variables

```
vi java8.sh
```

- Type the below under java8.sh file

```
export JAVA_HOME=/usr/lib/jvm/java-1.8.0-openjdk-  
amd64/  
export PATH=$PATH:$JAVA_HOME
```

- Bash the file

```
bash java8.sh
```

- Export java path

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
```

- Setting up a Non-Root User for Hadoop Environment. Therefore, Install the OpenSSH server and client using the following command:

```
sudo apt install openssh-server openssh-client -y
```

- Generate an SSH key pair and define the location it is to be stored in:

```
ssh-keygen -t rsa -P "" -f ~/.ssh/id_rsa
```

- Use the cat command to store the public key as authorized_keys in the ssh directory:

```
cat ~/.ssh/id_rsa.pub >> ~/.ssh/authorized_keys
```

- Set the permissions for your user with the chmod command:

```
chmod 0600 ~/.ssh/authorized_keys
```

- The new user is now able to SSH without needing to enter a password every time. Verify everything is set up correctly by using the hdoop user to SSH to localhost:

```
ssh localhost
```

- Open and Edit the .bashrc shell configuration file

```
sudo vi .bashrc
```

- Once you add the variables, save and exit the .bashrc file. It is vital to apply the changes to the current running environment by using the following command :-

#Hadoop Related Options

```
export HADOOP_HOME=/home/msrit/Downloads/hadoop-3.2.2 export HADOOP_INSTALL=$HADOOP_HOME export HADOOP_MAPRED_HOME=$HADOOP_HOME export HADOOP_COMMON_HOME=$HADOOP_HOME export HADOOP_HDFS_HOME=$HADOOP_HOME export YARN_HOME=$HADOOP_HOME export HADOOP_COMMON_LIB_NATIVE_DIR=$HADOOP_HOME/lib/native export PATH=$PATH:$HADOOP_HOME/sbin:$HADOOP_HOME/bin export HADOOP_OPTS="-Djava.library.path=$HADOOP_HOME/lib/native" export HIVE_HOME=/home/msrit/Downloads/apache-hive-3.1.2-bin export PATH=$PATH:$HIVE_HOME/bin
```

```
source ~/.bashrc
```

- Open the hadoop-env.sh file and make few changes:

```
sudo vi $HADOOP_HOME/etc/hadoop/hadoop-env.sh
```

- Uncomment the \$JAVA_HOME variable (i.e., remove the # sign) and add the full path to the OpenJDK installation on your system.

```
export JAVA_HOME=/usr/lib/jvm/java-8-openjdk-amd64
```

- Open the core-site.xml file

```
sudo vi $HADOOP_HOME/etc/hadoop/core-site.xml
```

- add the following configuration in between <configuration> and </configuration> to override the default values for the temporary directory and add your HDFS URL to replace the default local file system setting:

```
<configuration>
<property>
  <name>hadoop.tmp.dir</name>
  <value>/home/msrit/Downloads/tmpdata</value>
</property>
<property>
  <name>fs.default.name</name>
  <value>hdfs://127.0.0.1:9000</value>
</property>
</configuration>
```

- Use the following command to open the hdfs-site.xml file for editing:

```
sudo vi $HADOOP_HOME/etc/hadoop/hdfs-site.xml
```

- Add the following configuration in between <configuration> and </configuration>

```
<property>
<name>dfs.data.dir</name>
<value>/home/msrit/Downloads/dfsdata/namenode</value>
</property>
<property>
<name>dfs.data.dir</name>
<value>/home/msrit/Downloads/dfsdata/datanode</value>
</property>
<property>
<name>dfs.replication</name>
<value>1</value>
</property>
```

- Use the following command to access the mapred-site.xml file and define MapReduce values:

```
sudo vi $HADOOP_HOME/etc/hadoop/mapred-site.xml
```

- Add the following configuration in between <configuration> and </configuration> to change the default MapReduce framework name value to yarn:

```
<property>
  <name>mapreduce.framework.name</name>
  <value>yarn</value>
</property>
```

- Open the yarn-site.xml file in a text editor:

```
sudo vi $HADOOP_HOME/etc/hadoop/yarn-site.xml
```

- Add the following configuration in between <configuration> and </configuration>

```
<property>
  <name>yarn.nodemanager.aux-services</name>
  <value>mapreduce_shuffle</value>
</property>
<property>
  <name>yarn.nodemanager.aux-services.mapreduce.shuffle.class</name>
  <value>org.apache.hadoop.mapred.ShuffleHandler</value>
</property>
<property>
  <name>yarn.resourcemanager.hostname</name>
  <value>127.0.0.1</value>
</property>
<property>
  <name>yarn.acl.enable</name>
  <value>0</value>
</property>
<property>
  <name>yarn.nodemanager.env-whitelist</name>
  <value>JAVA_HOME,HADOOP_COMMON_HOME,HADOOP_HDFS_HOME,HADOOP_CONF_
_DIR,CLASSPATH_PERPEND_DISTCACHE,HADOOP_YARN_HOME,HADOOP_MAPRED_H
OME</value>
</property>
```

- Format the NameNode before starting Hadoop services for the first time

```
hdfs namenode -format
```

- cd to sbin directory of hadoop

```
cd Downloads/hadoop-3.2.2/sbin
```

- execute the following commands to start the NameNode and DataNode:

```
./start-dfs.sh
```

- Once the namenode, datanodes, and secondary namenode are up and running, start the YARN resource and nodemanagers by typing:

```
./start-yarn.sh
```

- Type this simple command to check if all the daemons are active and running as Java processes:

```
jps
```

- cd to the main directory

```
cd ..  
cd ..  
cd ..
```

Hive Installation

- Open terminal and cd into Downloads folder.
- Download Hive source code tarball using the following command in the terminal.

```
wget https://downloads.apache.org/hive/hive-3.1.2/apache-hive-3.1.2-bin.tar.gz
```

Or download the source code tarball directly from the [website](#).

- Unzip the file and go to home directory.

```
tar xzf apache-hive-3.1.2-bin.tar.gz
```

```
cd ..
```

- Set Hive globally by updating the system bash file.

```
gedit ~/.bashrc
```

- Paste these 2 lines at the end of the file

```
export HIVE_HOME=/home/msrit/Downloads/apache-hive-3.1.2-bin
```

```
export PATH=$PATH:$HIVE_HOME/bin
```

- Save and close the bashrc file and run the source command to load and save the new variables globally.

```
source ~/.bashrc
```


- Access the *hive-config.sh* file using the previously created **\$HIVE_HOME** variable:

```
sudo gedit $HIVE_HOME/bin/hive-config.sh
```

- Add the **HADOOP_HOME** variable and the full path to your Hadoop directory in *hive-config.sh* file. Save the edits and exit the *hive-config.sh* file.

```
export HADOOP_HOME=/home/msrit/hadoop-3.2.1
```

- Create a *tmp* directory within the HDFS storage layer. This directory is going to store the intermediary data Hive sends to the HDFS:

```
hdfs dfs -mkdir /tmp
```

- Add write and execute permissions to tmp group members:

```
hdfs dfs -chmod g+w /tmp
```

- Check if the permissions were added correctly:

```
hdfs dfs -ls /
```

- Create the *warehouse* directory within the */user/hive/* parent directory:

```
hdfs dfs -mkdir -p /user/hive/warehouse
```

- Add **write** and **execute** permissions to *warehouse* group members:

```
hdfs dfs -chmod g+w /user/hive/warehouse
```

- Check if the permissions were added correctly:

```
hdfs dfs -ls /user/hive
```

Apache Hive distributions contain template configuration files by default. The template files are located within the Hive *conf* directory and outline default Hive settings.

- Use the following command to locate the correct file:

```
cd $HIVE_HOME/conf
```

- Use the *hive-default.xml.template* to create the *hive-site.xml* file:

```
cp hive-default.xml.template hive-site.xml
```

- Access the *hive-site.xml* file using the nano text editor:

```
sudo gedit hive-site.xml
```

- You can configure the system to use your local storage rather than the HDFS layer by setting the *hive.metastore.warehouse.dir* parameter value to the location of your Hive *warehouse* directory.

Check the below lines are there in *hive-site.xml* file.

```
<property>
<name>hive.metastore.warehouse.dir</name>

<value>/user/hive/warehouse</value>
<description>location of default database for the warehouse
</description>
</property>
```

- Paste these lines to *hive-site.xml*:

```
<property>
<name>system:java.io.tmpdir</name>
<value>/tmp/hive/java</value>
</property>
<property>
<name>system:user.name</name>
<value>${user.name}</value>
</property>
<property>
<name>javax.jdo.option.ConnectionURL</name>
<value>jdbc:derby:/opt/hive-3.1.2-
bin/metastore_db;databaseName=metastore_db;create=true</value>
</property>
```

- Change the lines in *hive-site.xml* to given below lines:

```
<property>
<name>hive.txn.xlock.iow</name>
<value>true</value>
<description>
Ensures commands with OVERWRITE (such as INSERT
OVERWRITE) acquire Exclusive locks for transactional tables. This
ensures that inserts (w/o overwrite) running concurrently
are not hidden by the INSERT OVERWRITE.
</description>
</property>
```

- Locate the **guava jar** file in the Hive *lib* directory:

```
ls $HIVE_HOME/lib
```

- Locate the **guava jar** file in the Hadoop *lib* directory as well:

```
ls $HADOOP_HOME/share/hadoop/hdfs/lib
```

- Remove the existing **guava** file from the Hive *lib* directory:

```
rm $HIVE_HOME/lib/guava-19.0.jar
```

- Copy the **guava** file from the Hadoop *lib* directory to the Hive *lib* directory:

```
cp $HADOOP_HOME/share/hadoop/hdfs/lib/guava-27.0-jre.jar  
$HIVE_HOME/lib/
```

- Run the given command:

```
rm -rf metastore_db
```

- Use the **schematool** command once again to initiate the Derby database:

```
$HIVE_HOME/bin/schematool -dbType derby -initSchema
```

Steps To Run Hive

- Start the Hive command-line interface using the following commands:

```
cd $HIVE_HOME/bin
```

```
hive
```